


# Accelerating MATLAB algorithms

**By Dr Jasmina Lazić**  
**Application Engineer**  
**MathWorks**

# Agenda



Accelerating MATLAB code  
with parallel computing



Scaling up: accessing  
more/better hardware



Big data with MATLAB:  
an overview

# Agenda



Accelerating MATLAB code  
with parallel computing



Scaling up: accessing  
more/better hardware




Big data with MATLAB:  
an overview

# Accelerating MATLAB code

- Profile and optimise your code
- Introduction to parallel computing tools
- Solving big technical problems in MATLAB
  - Programming task parallel applications
  - Programming data parallel applications

# Who wants faster code?



# Using More Hardware

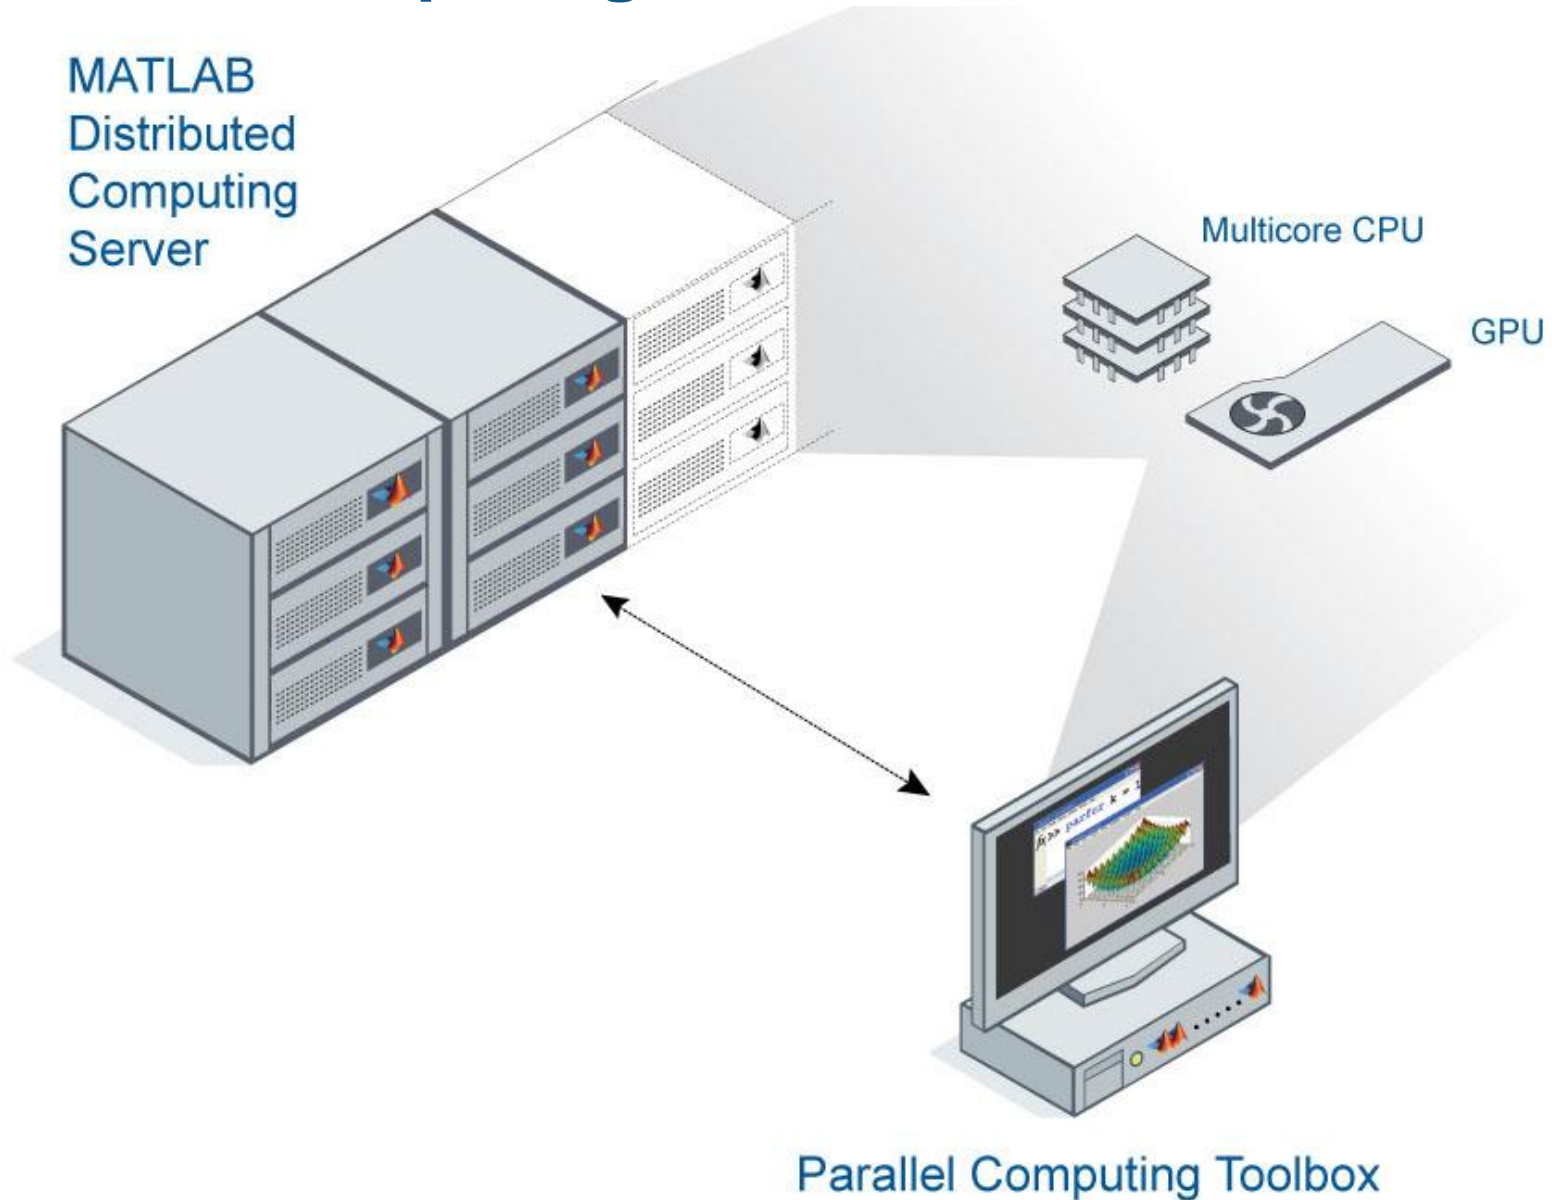
## Built-in multithreading

- Automatically enabled in MATLAB since R2008a
- Multiple threads in a single MATLAB computation engine  
[www.mathworks.com/discovery/multicore-matlab.html](http://www.mathworks.com/discovery/multicore-matlab.html)

## Parallel computing using explicit techniques

- Multiple computation engines controlled by a single session
- High-level constructs to let you parallelise MATLAB applications
- Perform MATLAB computations on GPUs

# Parallel Computing Products



# Solving Big Technical Problems

## Challenges

You could...

## Solutions

Long running

---

Computationally  
intensive

Wait



Larger Compute Pool  
(e.g. More Processors)

---

Large data set

Reduce size  
of problem



Larger Memory Pool  
(e.g. More Machines)






## Optimizing JIT Steel Manufacturing Schedule

Cut simulation time from 1 hour to 5 minutes

## Heart Transplant Studies

3-4 weeks reduced to 5 days




## Flight Test Data Analysis

16x Faster

## Mobile Communications Technology

Simulation time reduced from weeks to hours,  
5x more scenarios



## Hedge Fund Portfolio Management

Simulation time reduced from 6 hours to 1.2 hours

# Programming Parallel Applications (CPU)



Ease of Use

- Parallel enabled functionality in toolboxes

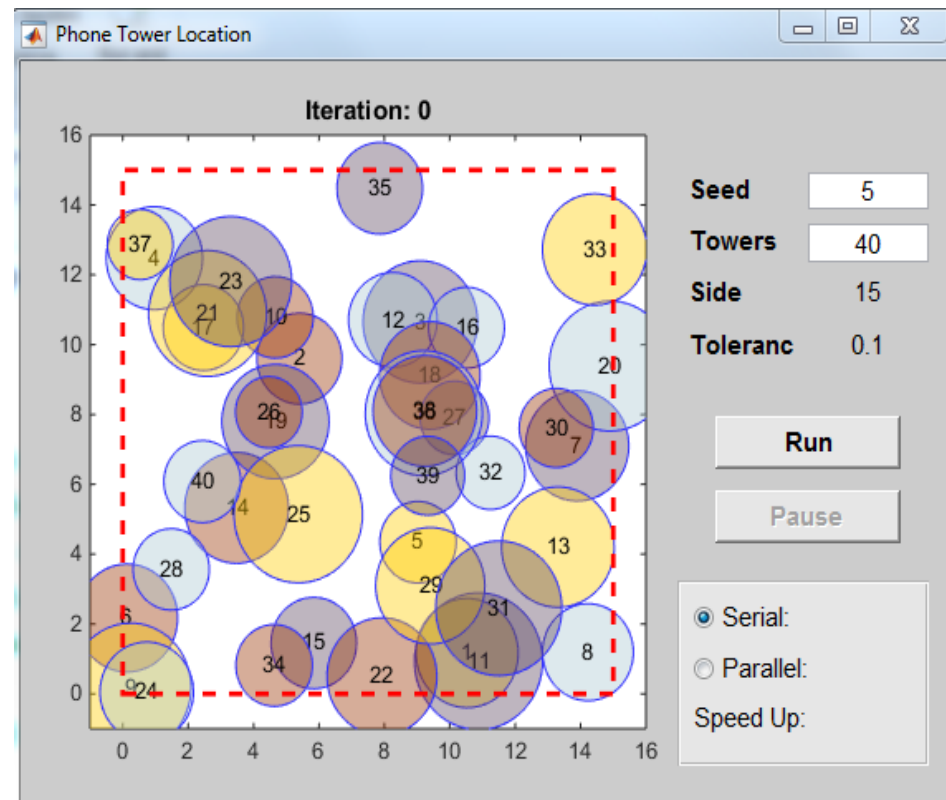


Greater Control

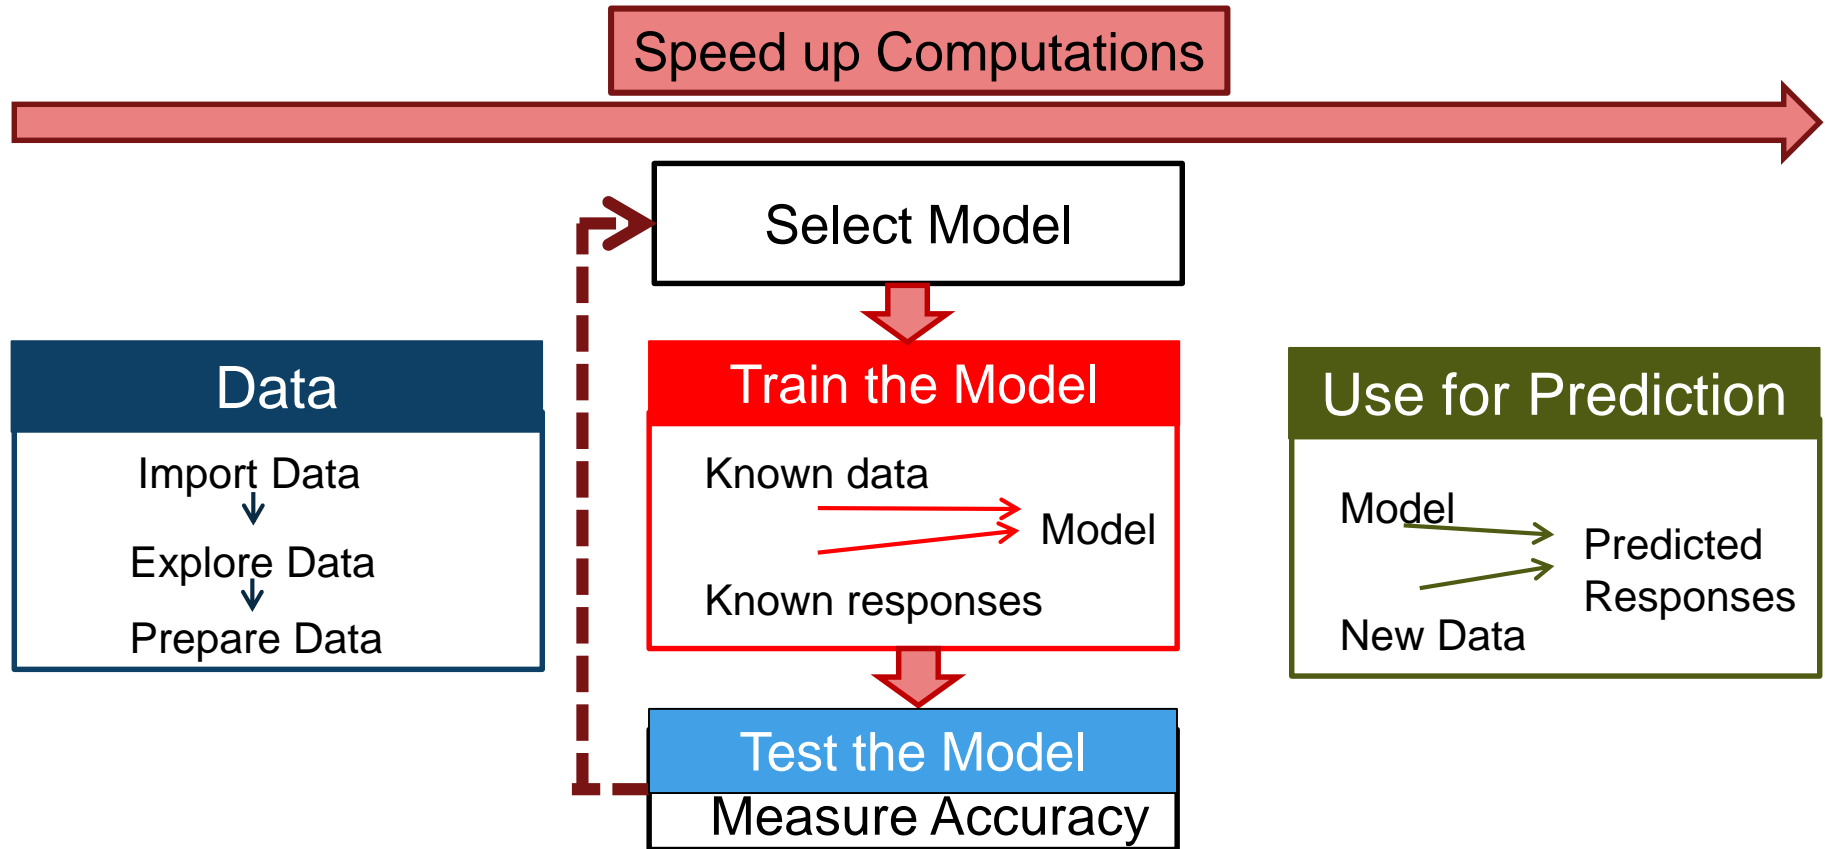
# Example: Optimizing Cell Tower Position

## Parallel enabled functionality

- With Parallel Computing Toolbox use parallel enabled algorithms in Optimization Toolbox
- Run optimization in parallel
- Use pool of MATLAB workers




# Supervised Learning - Workflow



# Tools Providing Parallel Computing Support

- Optimization Toolbox
- Global Optimization Toolbox
- Statistics Toolbox and MachineLearning Toolbox
- Signal Processing Toolbox
- Neural Network Toolbox
- Image Processing Toolbox



**'UseParallel', true**  
**options = optimset('UseParallel', true);**

*Directly leverage functions in Parallel Computing Toolbox*

[www.mathworks.com/builtin-parallel-support](http://www.mathworks.com/builtin-parallel-support)

# Programming Parallel Applications (CPU)



Ease of Use

- Parallel enabled functionality in toolboxes
- Simple programming constructs:  
`parfor`, `batch`, `distributed`



Greater Control

# Example: Parameter Sweep of ODEs

## Parallel for-loops



- Parameter sweep of ODE system

- Damped spring oscillator
- Sweep through different values of damping and stiffness
- Record peak value for each simulation

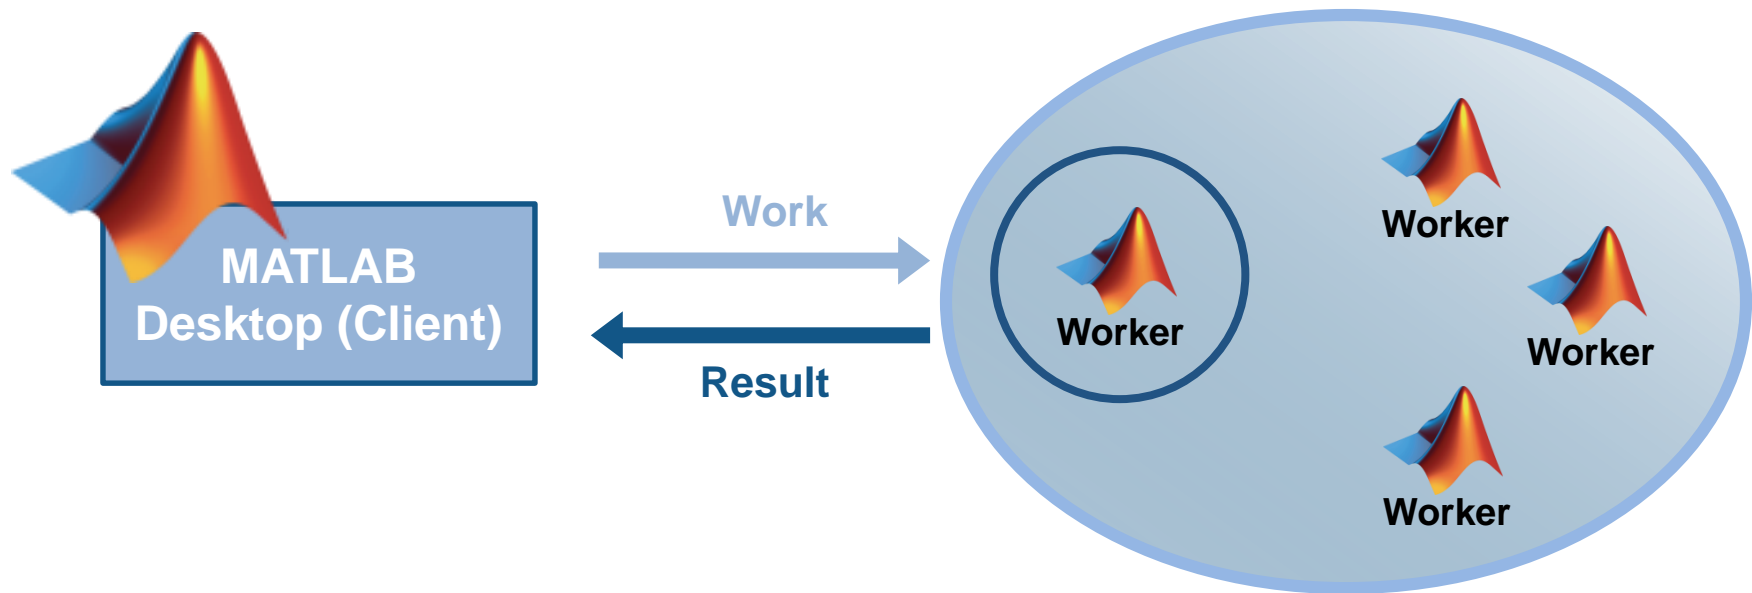
- Convert `for` to `parfor`

- Use pool of MATLAB workers

$$\overset{5}{\underbrace{m}} \ddot{x} + \underbrace{b}_{1,2,\dots} \dot{x} + \underbrace{k}_{1,2,\dots} x = 0$$



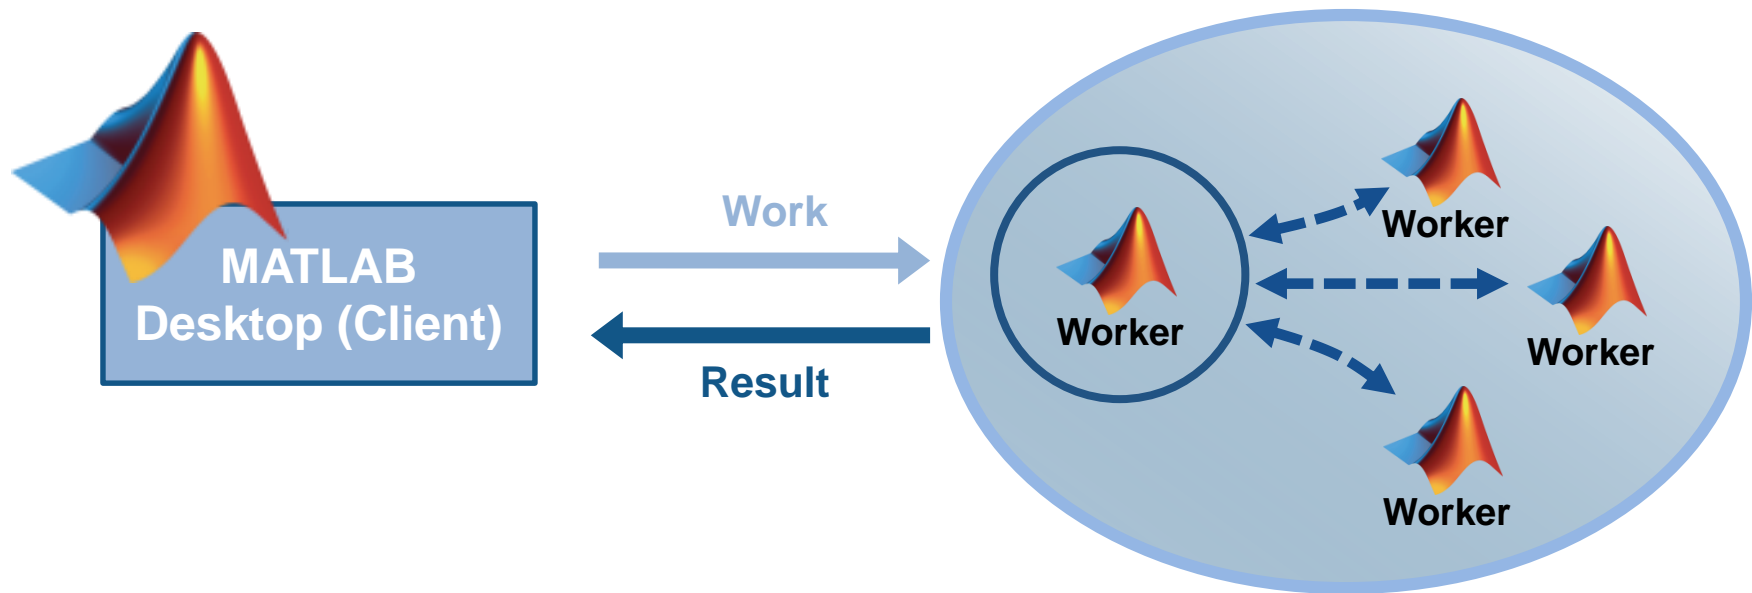
# Offload Computations with batch



`batch (...)`



# Offload and Scale Computations with `batch` & `pool`




`batch(..., 'Pool', ...)`

# distributed

Split data over multiple cores

```
parpool('local')  
A = distributed.randn(1000);  
[V,D] = eig((A+A')/2);  
D = gather(D);
```



# Programming Parallel Applications (CPU)



Ease of Use

- Parallel enabled functionality in toolboxes
- Simple programming constructs: `parfor`, `batch`, `distributed`
- Advanced programming constructs: `createJob`, `spmd`, `labSend/labReceive`



Greater Control

# Agenda



Accelerating MATLAB code  
with parallel computing



Scaling up: accessing  
more/better hardware




Big data with MATLAB:  
an overview

# Using more hardware: why and when?



We have seen that with PCT we get some speed up locally and we have access to a powerful server or cluster with MDCS



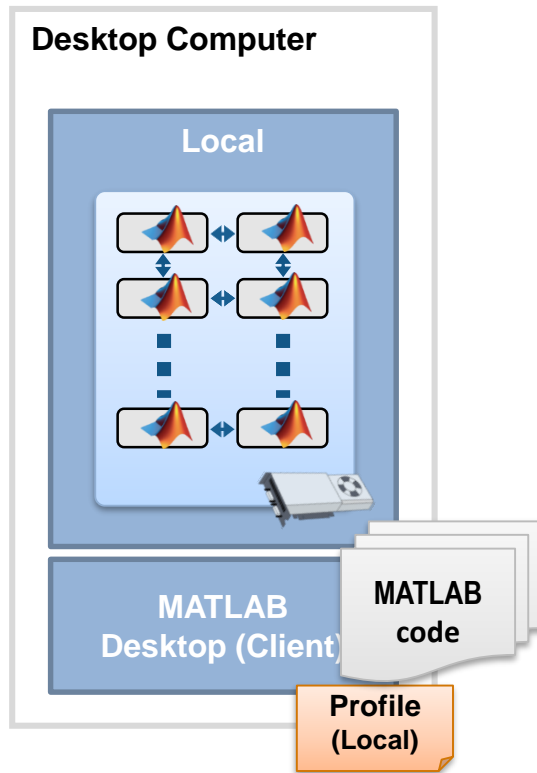
Our data is too large to fit in the desktop memory



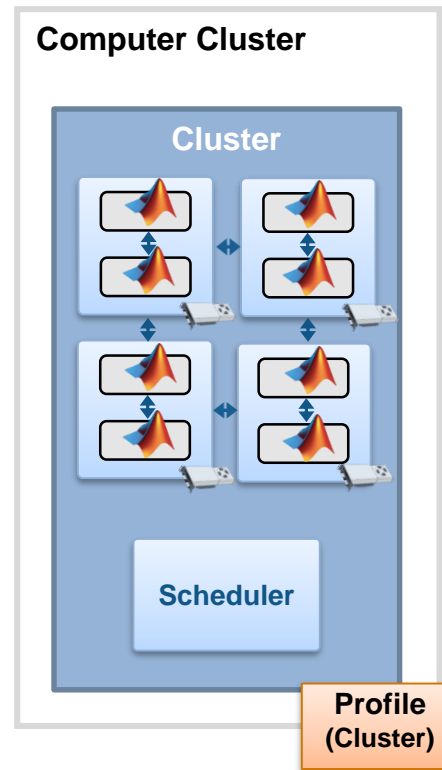
We have access to GPUs

# Use MATLAB Distributed Computing Server

## 1. Prototype code

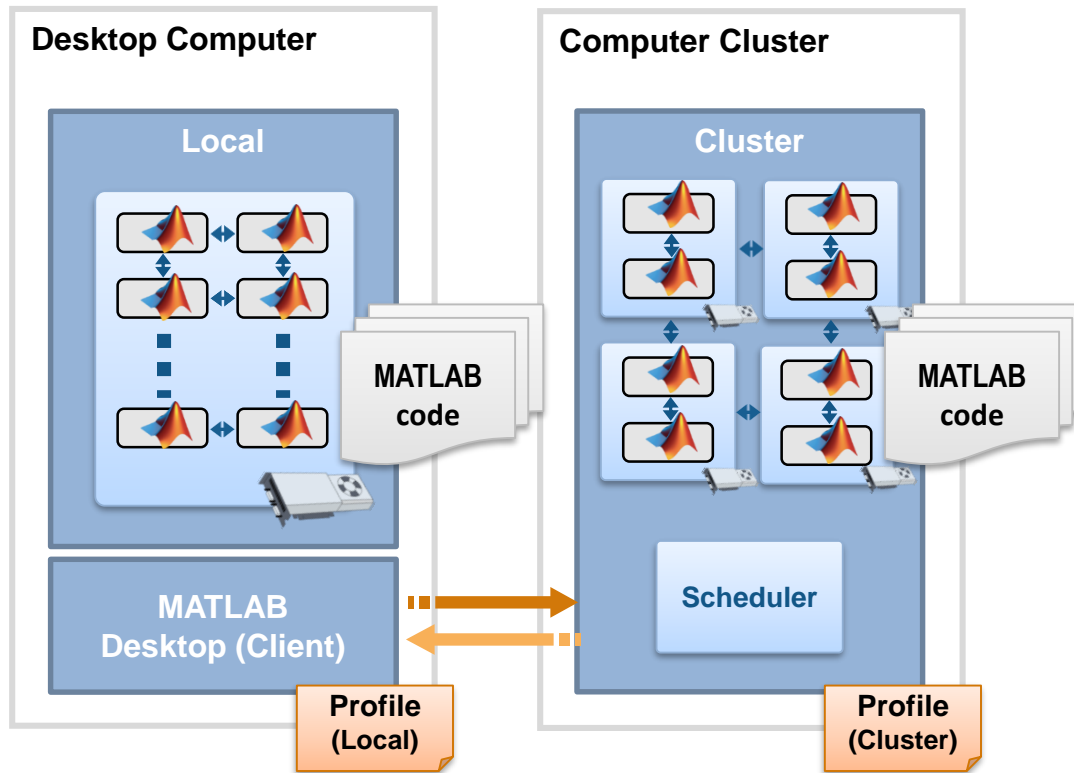


# Use MATLAB Distributed Computing Server



1. Prototype code
2. Get access to an enabled cluster

# Use MATLAB Distributed Computing Server





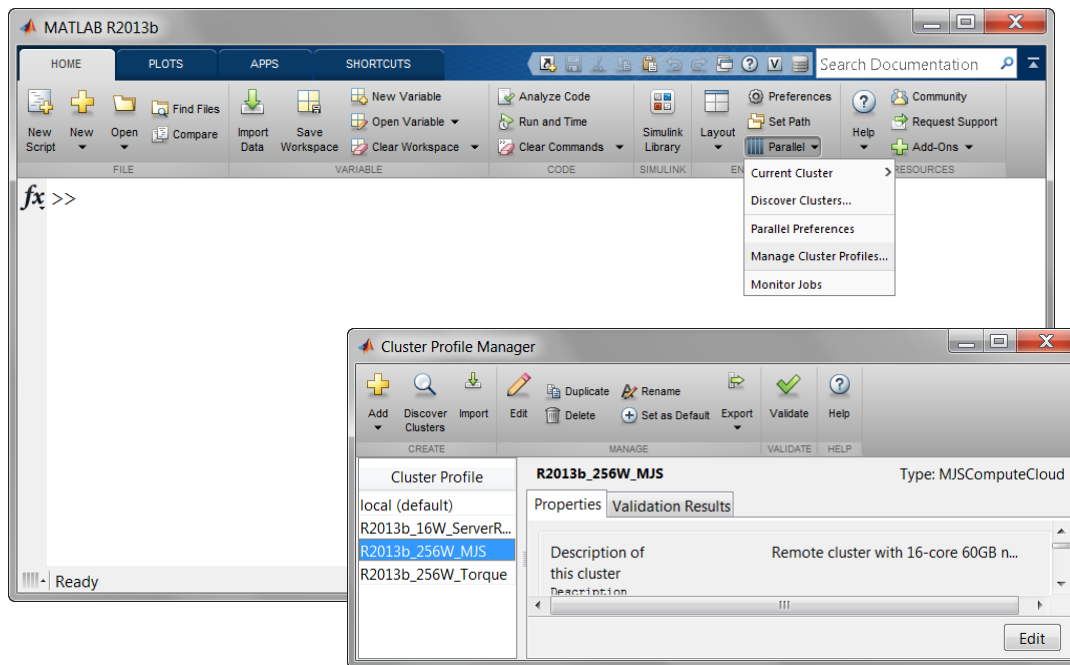
1. Prototype code
2. Get access to an enabled cluster
3. Switch cluster profile to run on cluster resources



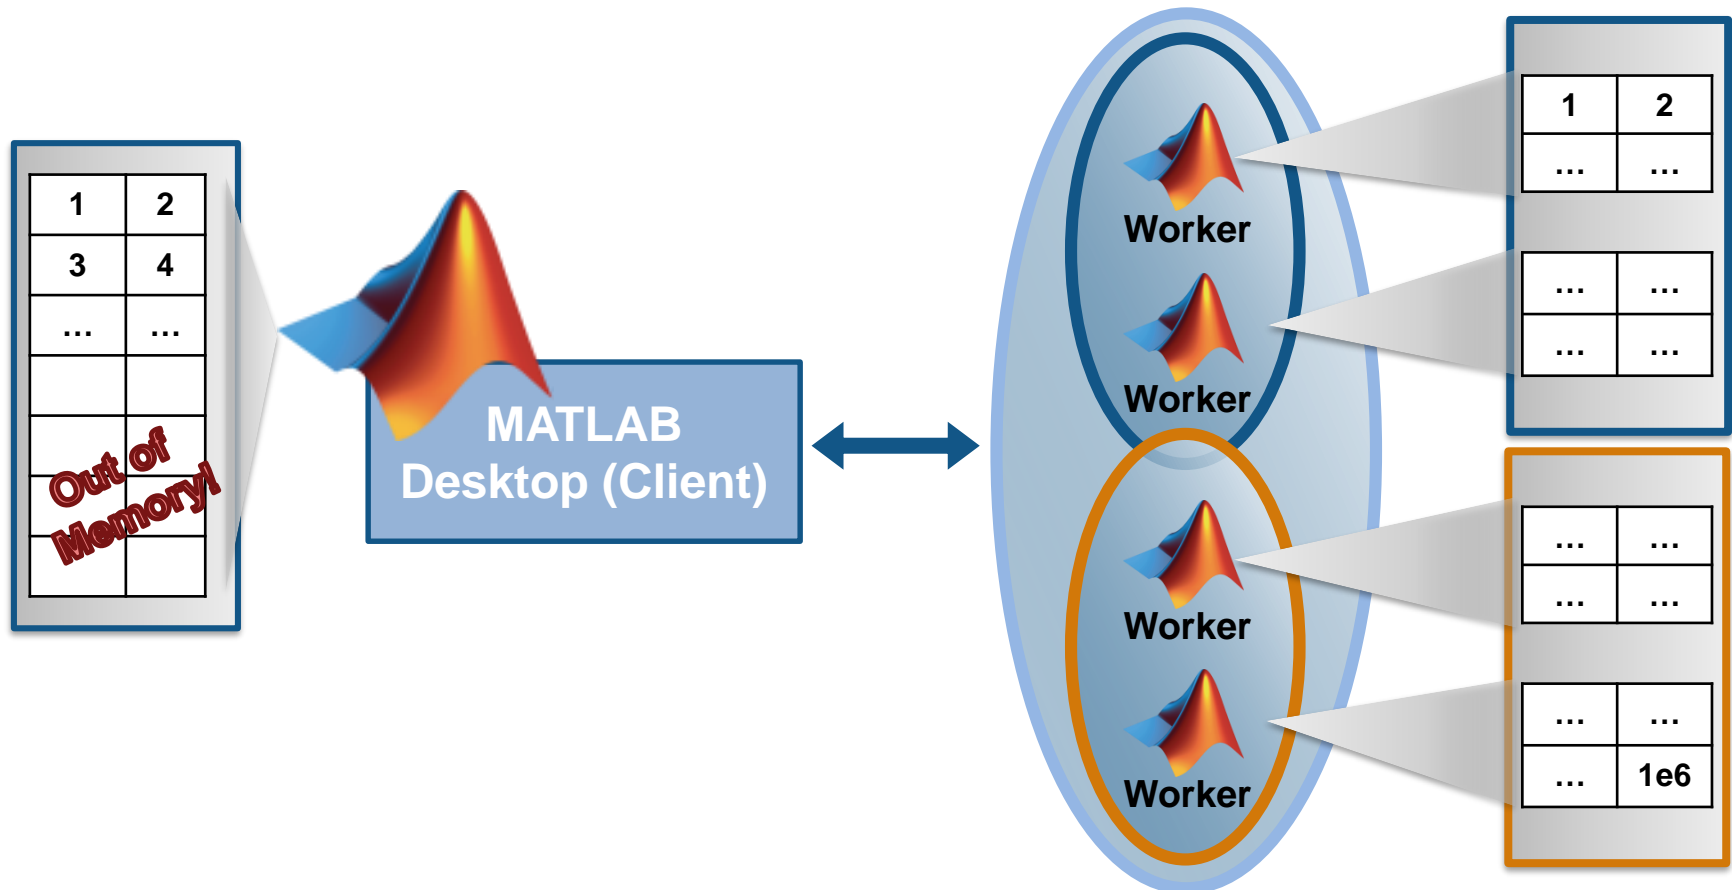
# Example: Migrate from Desktop to Cluster

- Change hardware without changing algorithmic code

$$\overset{5}{m}\ddot{x} + \underset{1,2,\dots}{b}\dot{x} + \underset{1,2,\dots}{k}x = 0$$




# Overcoming Desktop Memory Limitations



# Distributed Memory Calculations

## Multiplication of 2 NxN matrices

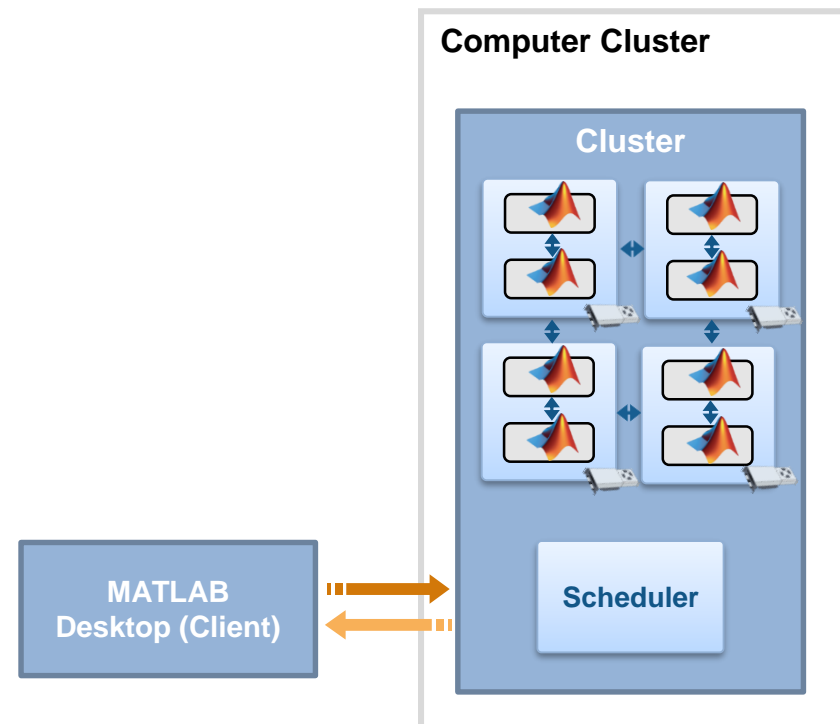


N	Execution time (seconds)		
	1 node, 16 workers	2 nodes, 32 workers	4 nodes, 64 workers
8000	19	13	11
16000	120	75	50
20000	225	132	86
25000	-	243	154
30000	-	406	248
35000	-	-	376
45000	-	-	743
50000	-	-	-

Processor: Intel Xeon E5-class v2  
16 cores, 60 GB RAM per compute node, 10 Gb Ethernet


# Take Advantage of Cluster Hardware

- Offload computation:
  - Free up desktop
  - Access better computers
- Scale speed-up:
  - Use more cores
  - Go from hours to minutes
- Scale memory:
  - Utilize distributed arrays
  - Solve larger problems without re-coding algorithms
- Leverage supplied infrastructure
  - File transfer / path augmentation
  - Job monitoring



# Further Scaling Big Data Capacity

MATLAB supports a number of programming constructs for use with clusters



- General compute clusters
  - *Parallel for loops – embarrassingly parallel algorithms*
  - *SPMD – distributed processing*
- Hadoop clusters
  - *MapReduce – analyze data stored in the Hadoop Distributed File System*

# What is a Graphics Processing Unit (GPU)

- Originally for graphics acceleration, now also used for scientific calculations
- Massively parallel array of integer and floating point processors
  - Typically hundreds of processors per card
  - GPU cores complement CPU cores
- Dedicated high-speed memory



\* Parallel Computing Toolbox requires NVIDIA GPUs with Compute Capability 1.3 or higher, including NVIDIA Tesla 20-series products. See a complete listing at [www.nvidia.com/object/cuda\\_gpus.html](http://www.nvidia.com/object/cuda_gpus.html)

# Programming Parallel Applications (GPU)



Ease of Use

- Built-in support with Toolboxes
- Simple programming constructs:  
`gpuArray`, `gather`
- Advanced programming constructs:  
`arrayfun`, `bsxfun`, `spmd`
- Interface for experts:  
`CUDAKernel`, `MEX` support



Greater Control

# Accessing GPUs (Graphics Processing Units)

```
>> A = someArray(200, 10000);  
>> G = gpuArray(A); % Push to GPU memory  
>> F = cov(G);  
>> B = gather(F) % Bring back into MATLAB
```



To access GPU functionality you will need R2010b or later, and NVIDIA GPUs with Compute Capability 1.3 or greater including NVIDIA Tesla 10-series and 20-series products: [www.nvidia.com/object/cuda\\_gpus.html](http://www.nvidia.com/object/cuda_gpus.html)



# CPU/GPU comparison

**%% CPU**

```
Acpu = rand(2000);
```

```
tic
```

```
Bcpu = cov(Acpu);
```

```
toc
```

Elapsed time is 0.264531 seconds.

**%% GPU**

```
Agpu = gpuArray.rand(2000);
```

```
tic
```

```
Bgpu = cov(Agpu);
```

```
wait(gpuDevice);
```

```
toc
```


```
B = gather(Bgpu);
```

Elapsed time is 0.041266 seconds.

# Summary

- Develop parallel MATLAB applications without being a parallel programming expert
- Speed up the execution of your MATLAB applications using additional hardware
- Prototype on your desktop and easily scale to a cluster

# Agenda



Accelerating MATLAB code  
with parallel computing



Scaling up: accessing  
more/better hardware



Big data with MATLAB:  
an overview

# Big Data

*“Any collection of data sets so large and complex that it becomes difficult to process using... traditional data processing applications.”*

*(Wikipedia)*

## FINANCE

Market Risk,  
Regulatory



## AUTO

Fleet Data  
Analysis



## Medical Devices

Patient Outcomes



## ENERGY

Asset  
Optimization



- Access to Data
- Development of scalable algorithms
- Rapid data exploration
- Ease of deployment

# Considerations: large data analytics

## *Data Characteristics*



1. Size & type of data?
2. Where is your data?
3. What hardware do you have access to?
4. Analysis Characteristics
  - Embarrassingly Parallel
  - Analyse sub-segments of data and aggregate results
  - Operate on entire dataset

# Big Data Capabilities in MATLAB

## Memory and Data Access


- 64-bit processors
- Memory Mapped Variables
- Disk Variables
- Databases
- Datastores

## Programming Constructs

- Streaming
- Block Processing
- Parallel-for loops
- GPU Arrays
- SPMD and Distributed Arrays
- MapReduce

## Platforms

- Desktop (Multicore, GPU)
- Clusters
- Cloud Computing (MDCS on EC2)
- Hadoop



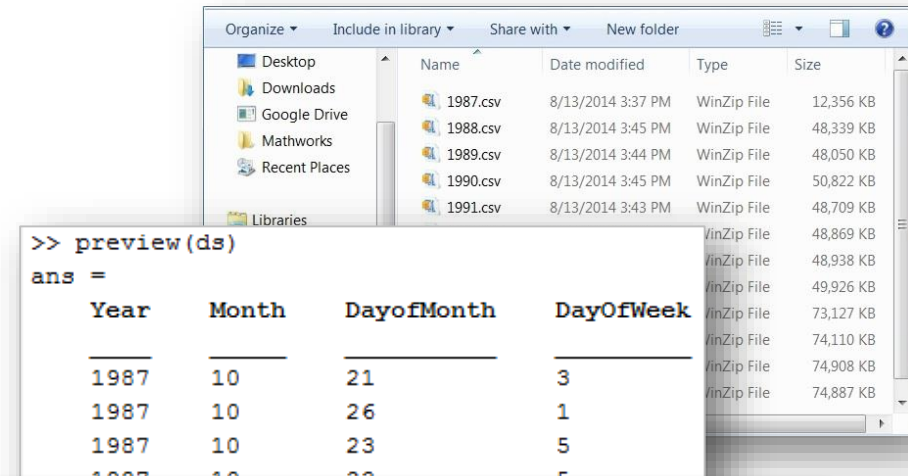
# Big Data on the Desktop

## MATLAB

### datastore

*Access files or collections of files that do not fit into memory*

- Preview data structure and format
- Import data using column names
- Incremental access to data



```
airdata = datastore('*.csv');
airdata.SelectedVariables = {'Distance', 'ArrDelay'};

data = read(airdata);
```


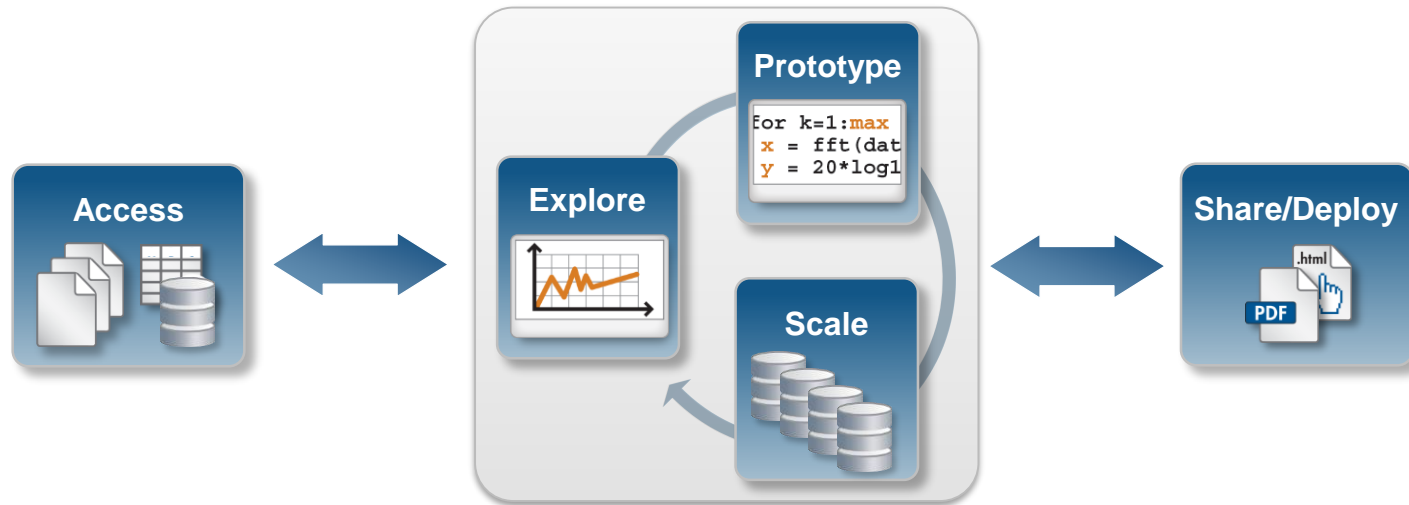
### mapreduce

*Programming technique for analyzing data sets that do not fit into memory*


- Tabular text files or big database tables
- Increase compute capacity with parallel processing
- Access HDFS to develop algorithms for use on Hadoop

```
*****
*           MAPREDUCE PROGRESS           *
*****
Map 0%      Reduce 0%
Map 20%     Reduce 0%
Map 40%     Reduce 0%
Map 60%     Reduce 0%
Map 80%     Reduce 0%
Map 100%    Reduce 25%
Map 100%    Reduce 50%
Map 100%    Reduce 75%
Map 100%    Reduce 100%
```

# Big Data Analytics with MATLAB



Work on the desktop



Scale capacity as needed